

MASTER IN
COMPUTER
SCIENCE



Contrastive Learning for Character Detection in Ancient Greek Papyri

Advisors: Prof. Rolf Ingold,
Prof. Andreas Fischer
Supervisor: Lars Vögtlin

Master Thesis
September 6th, 2024

Vedasri Nakka
Vedasri.nakka@unine.ch





Agenda

Introduction

Methods

Results

Discussion & limitations

Conclusion & Future work

Research Question

How effective is SimCLR contrastive learning for Greek-letter recognition tasks compared to traditional baseline methods, and which data augmentation strategies enhance its performance?

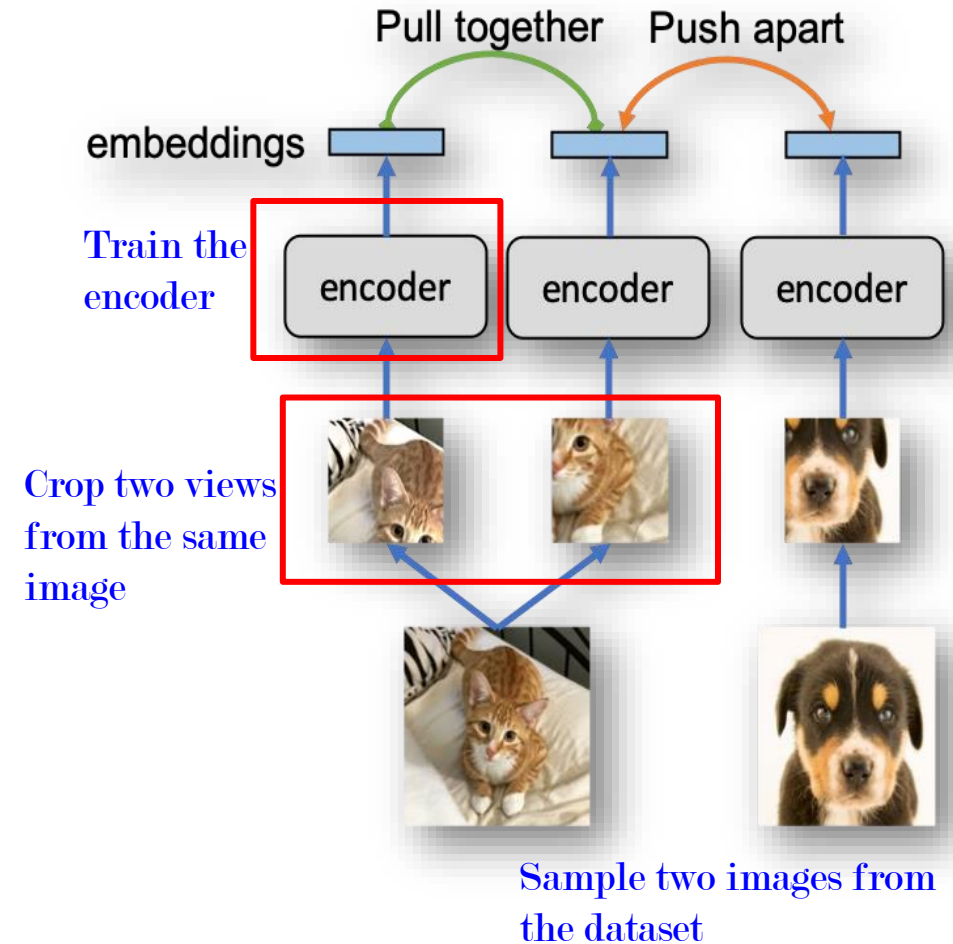
Overview of Contrastive Learning

Contrastive Learning is a **Self-Supervised Learning** technique, that teaches the model to learn embeddings of unlabeled data by

- **maximizing** the similarity between similar instances and ;
- **minimizing** the similarity between dissimilar instances.

Major Benefit:

- CL do not require **labels** to conduct the training and thus can be pretrained on large unlabeled corpus to obtain a powerful feature encoder.



Training Strategies

Baseline Model

Supervised learning model used for classification tasks utilizing cross-Entropy Loss.

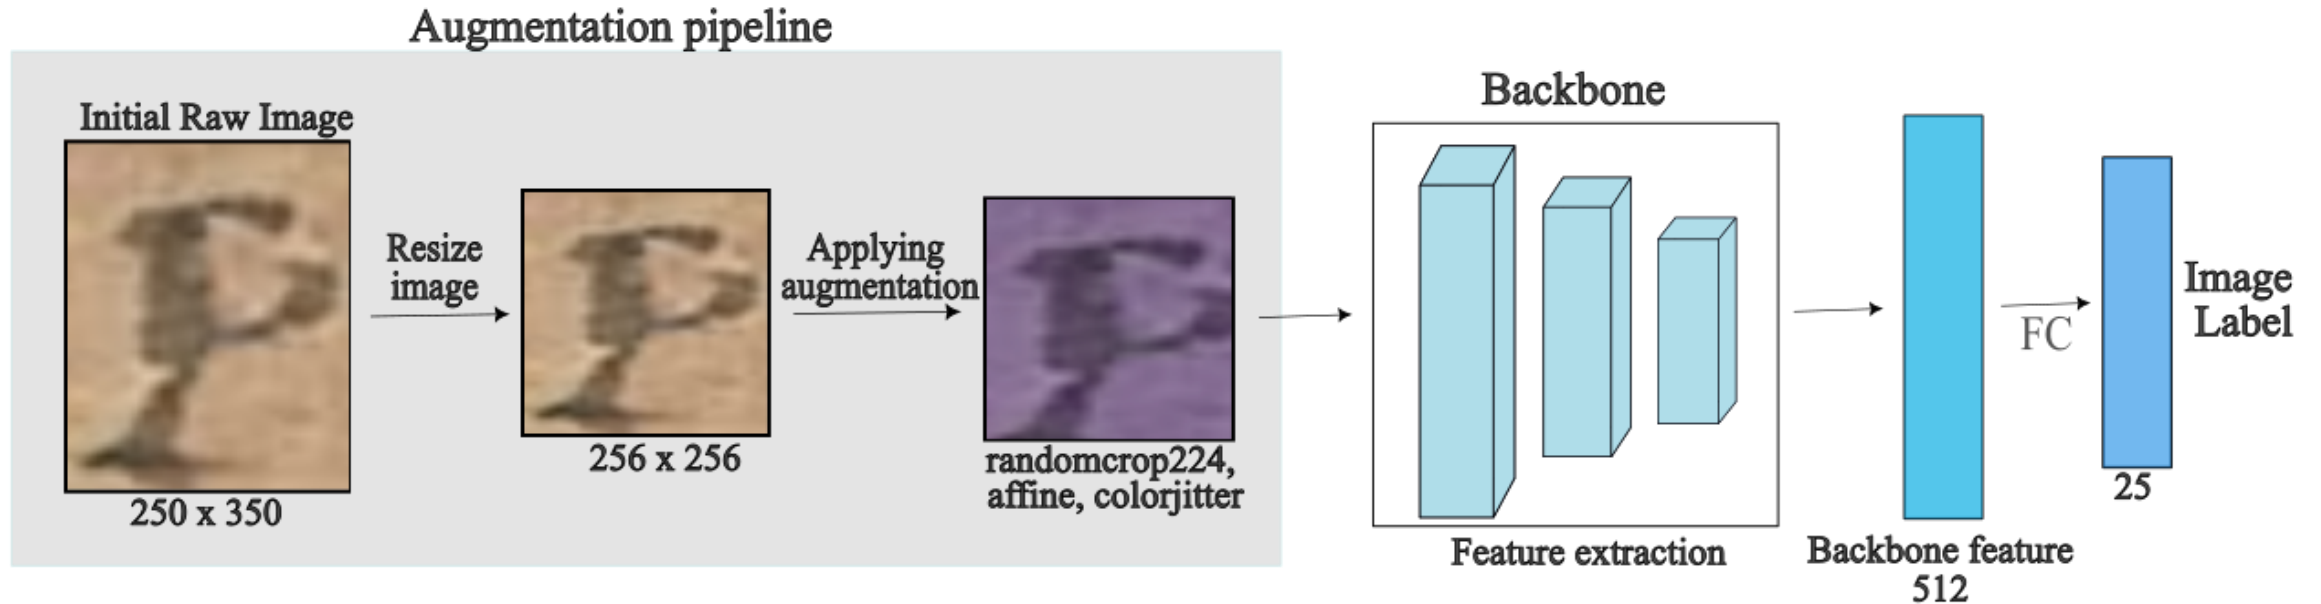
Triplet Embedding Model

Minimize the distance between the anchor and positive image embeddings while maximizing the distance between anchor with the negative image embeddings

SimCLR Model

A self-supervised learning method using positive pairs (augmented views of the same image) and contrasts them with negative pairs (different images) to learn feature representations using Contrastive Loss.

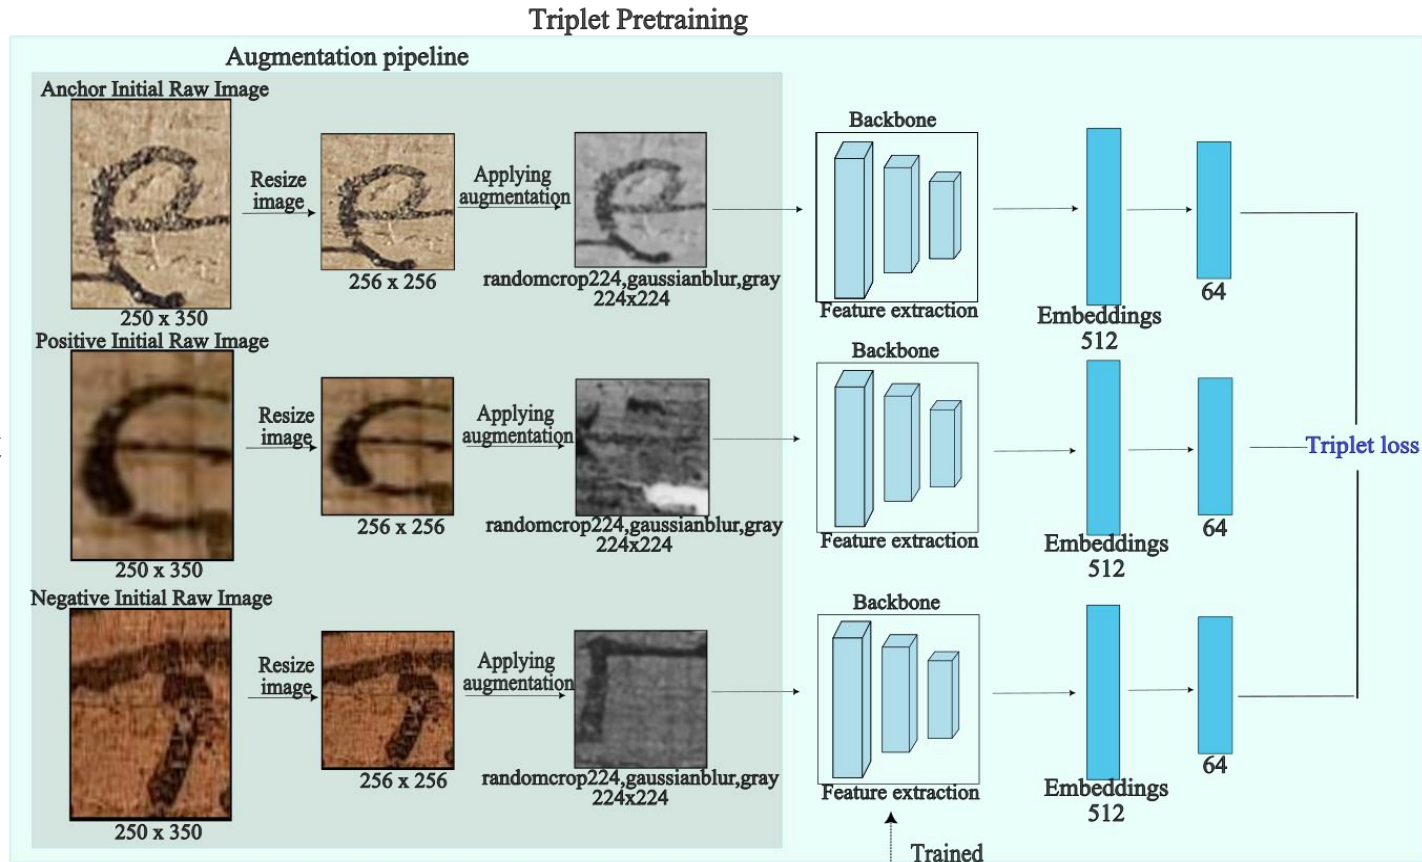
Baseline model pipeline



- We first perform the spatial and pixel-level augmentation to obtain the final pre-processed image. We send the augmented image to the model which extracts a feature vector at the end of the backbone. The final feature vector is sent to the classification layer to get output probability.

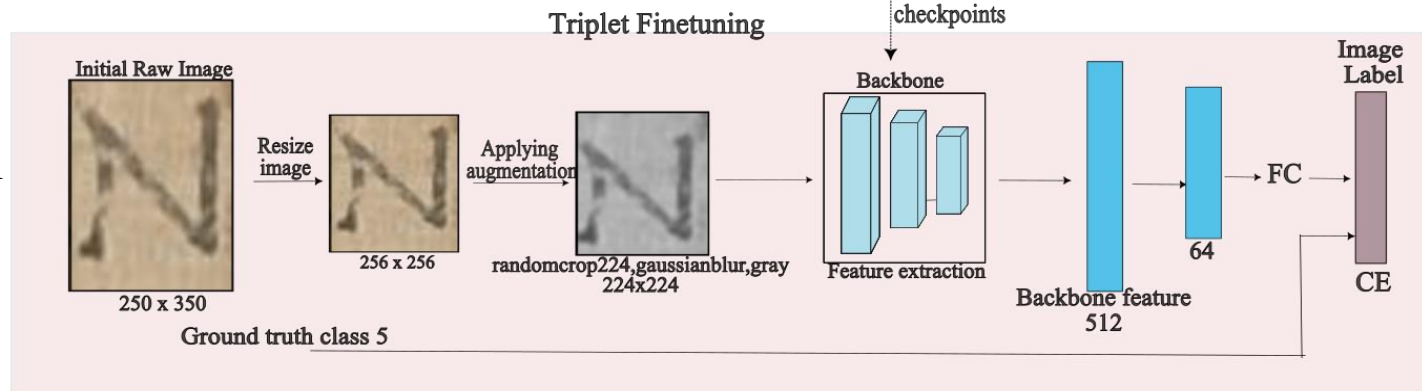
Triplet model pipeline

Alpub
dataset

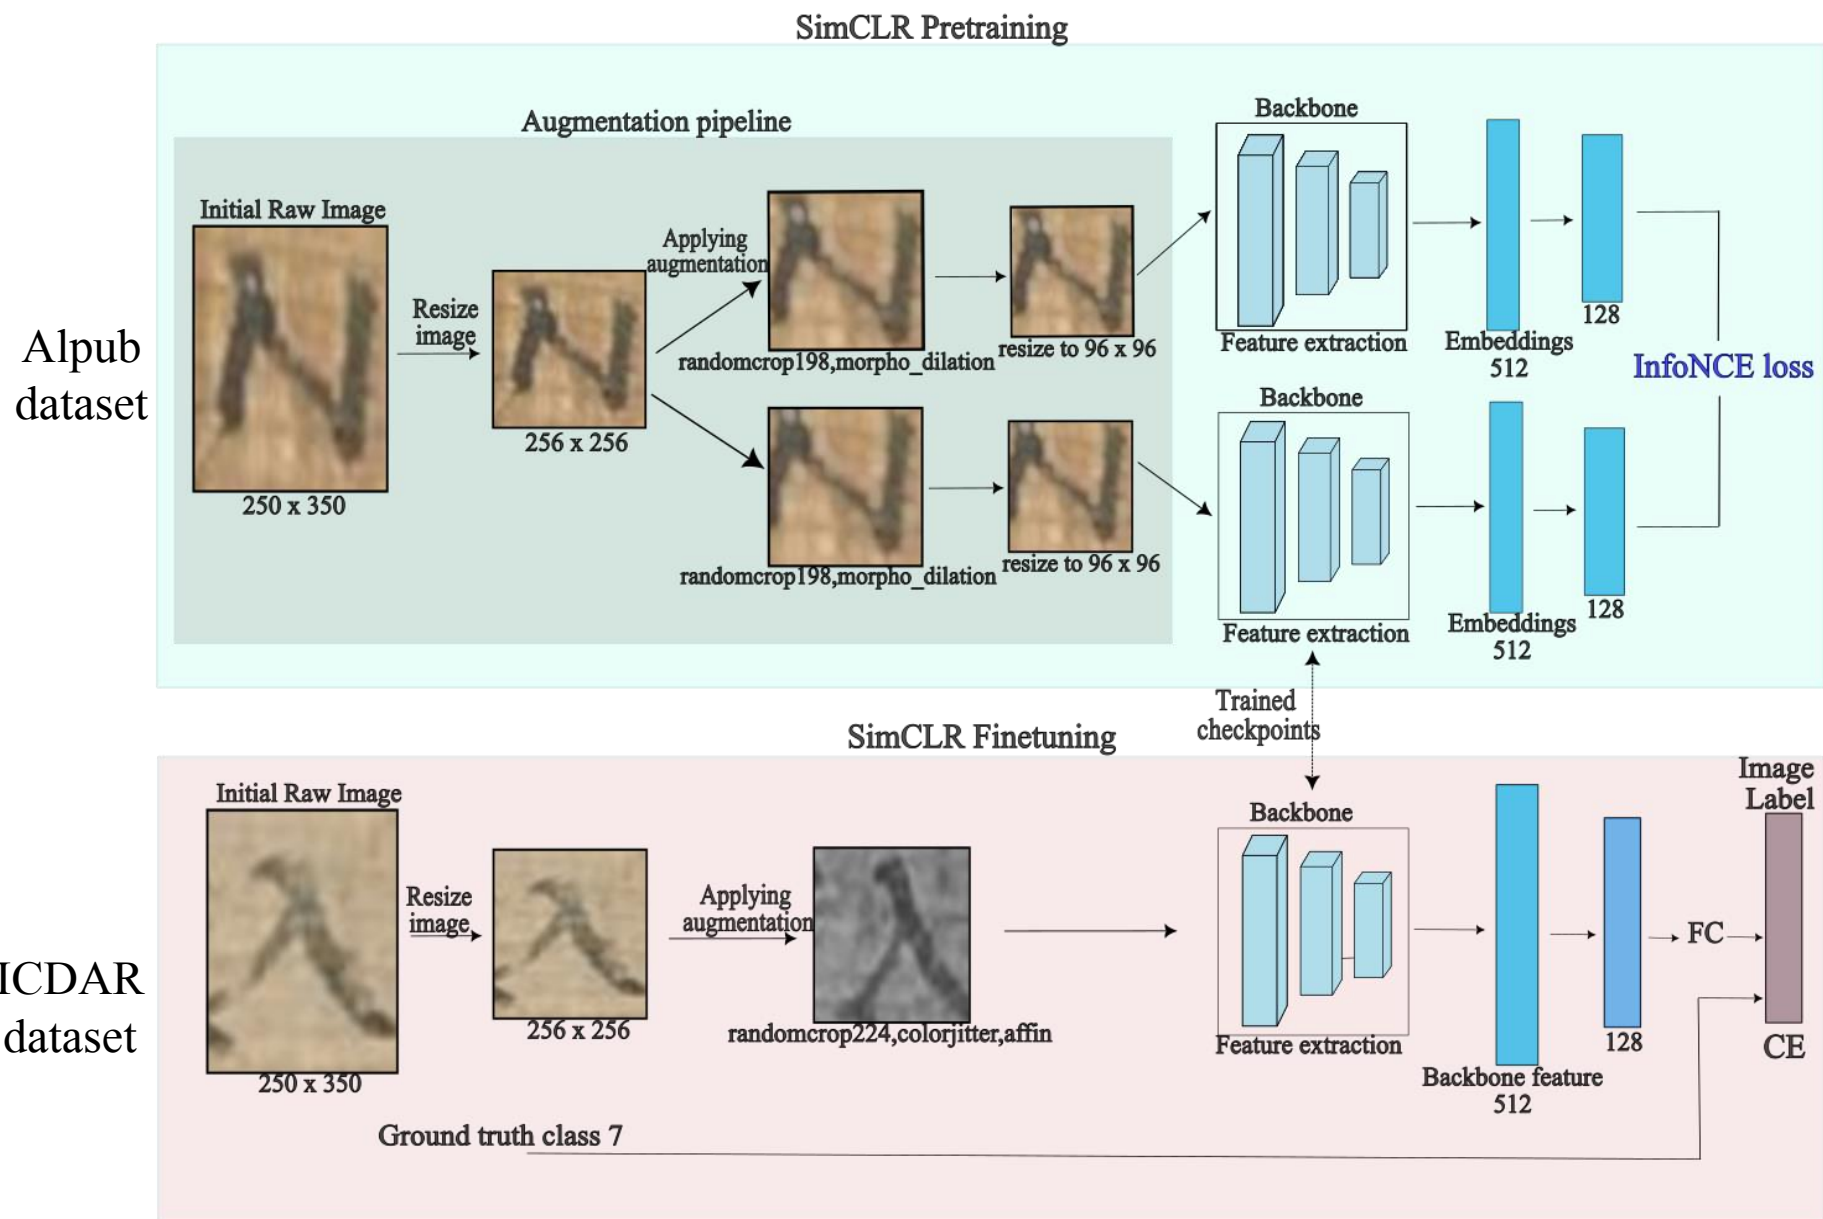


- The top block, highlighted with a green background, represents the pretraining stage, where the model is trained on the pretraining dataset using triplet loss to learn embeddings from triplet pairs.
- In the next stage, a classification layer is added on top of the embedding layer, and the model is trained end-to-end with cross-entropy (CE) loss.

ICDAR
dataset



SimCLR model pipeline



- The top block, highlighted with a green background, represents the pretraining stage, where the model is trained on the pretraining dataset using InfoNCE loss to learn embeddings from SimCLR augmented view of same image.
- In the next stage, a classification layer is added on top of the embedding layer, and the model is trained end-to-end with cross-entropy (CE) loss.

Comparison of batch data for 3 Models

Baseline model



Batch of data from 16 datapoints – 4 order


Triplet model



The dataloader selects 8 images from the same class as the anchor images (positives) and 8 images from different classes (negatives),

Batch of data from 8 datapoints (each datapoint contributes to anchor, positive, negative pairs)

SimCLR model



Pair of images (top and bottom) cropped from the same image with augmentation

Batch of data from 8 datapoints (each datapoint contributes to pair of images)

Datasets

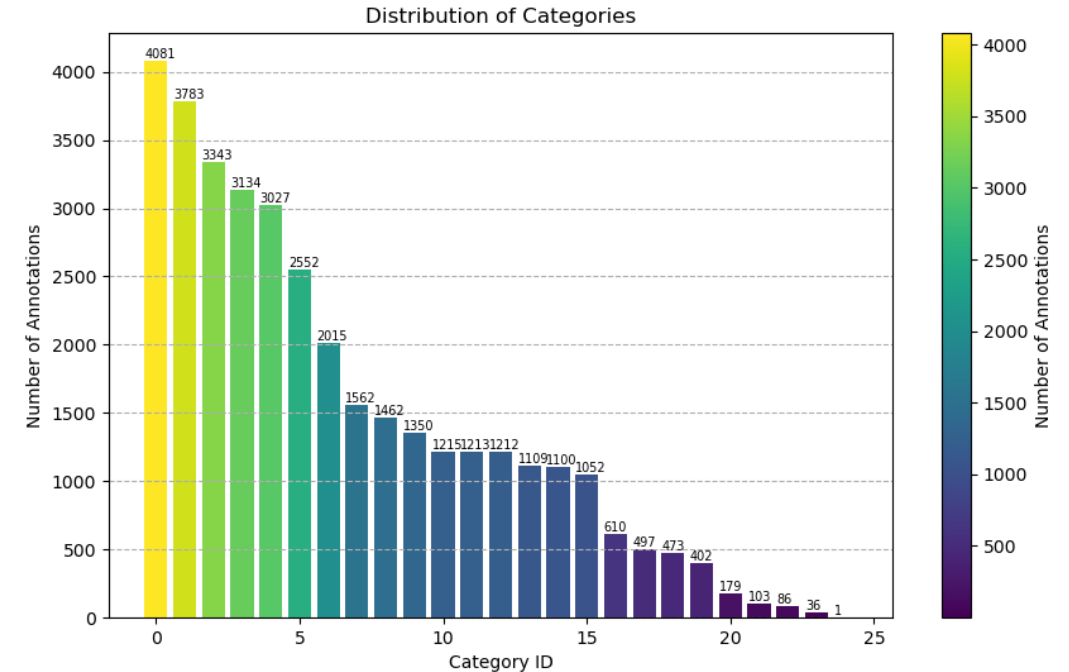
Pre-training dataset:

- AL-PUB dataset containing 205,797 cropped images
- Total of 24 Greek Letter classes

Fine-tuning dataset:

The 'ICDAR 2023 Competition on Detection and Recognition of Greek Letters on Papyri'

- Contains 153 training images and 34 test images of full document-level image resolution
- On cropping each Greek character from full image in dataset, we obtain around 34K cropped Greek Letter images in the training set
- Total of 25 Greek Letter classes



Cropped Greek Letters



Albumentation Library

- Albumentations is a Python library for image augmentation.
- Albumentations supports all common computer vision tasks such as classification, semantic segmentation, instance segmentation, object detection, and pose estimation.
- The library contains more than 70 different augmentations to generate new training samples from the existing data.

Types of Augmentations:

1. Pixel-level transforms: These transformations modify individual pixels in an image.

2. Spatial-level transforms: These transformations change the overall structure or geometry of the image.



List of augmentations

Pixel-level transforms

Pixel-level transforms will change just an input image and will leave any a boxes, and keypoints unchanged. The list of pixel-level transforms:

- [AdvancedBlur](#)
- [Blur](#)
- [CLAHE](#)
- [ChannelDropout](#)
- [ChannelShuffle](#)
- [ChromaticAberration](#)
- [ColorJitter](#)
- [Defocus](#)
- [Downscale](#)
- [Emboss](#)
- [Equalize](#)
- [FDA](#)
- [FancyPCA](#)
- [FromFloat](#)
- [GaussNoise](#)
- [GaussianBlur](#)
- [GlassBlur](#)
- [HistogramMatching](#)
- [HueSaturationValue](#)
- [ISONoise](#)
- [ImageCompression](#)

Spatial-level transforms

Spatial-level transforms will simultaneously change both an input image as well bounding boxes, and keypoints. The following table shows which additional tar

Transform	Image	Mask	BBoxes	Keypoints
Affine	✓	✓	✓	✓
BBoxSafeRandomCrop	✓	✓	✓	
CenterCrop	✓	✓	✓	✓
CoarseDropout	✓	✓		✓
Crop	✓	✓	✓	✓
CropAndPad	✓	✓	✓	✓
CropNonEmptyMaskIfExists	✓	✓	✓	✓
D4	✓	✓	✓	✓
ElasticTransform	✓	✓	✓	
Flip	✓	✓	✓	✓
GridDistortion	✓	✓	✓	
GridDropout	✓	✓		
HorizontalFlip	✓	✓	✓	✓
Lambda	✓	✓	✓	✓
LongestMaxSize	✓	✓	✓	✓
MaskDropout	✓	✓		
Morphological	✓	✓		
NoOp	✓	✓	✓	✓

Data augmentation dictionary

```
transform_dict = {
    "default": A.Resize(height=256, width=256, p=1.0),
    "resize256": A.Resize(height=256, width=256, p=1),
    "randomcrop224": A.RandomCrop(224, 224, p=1.0),
    "randomcrop198": A.RandomCrop(198, 198, p=1.0),
    "randomcrop128": A.RandomCrop(128, 128, p=1.0),
    "randomcrop180": A.RandomCrop(180, 180, p=1.0),

    # scale hyperparam
    "morpho_dilation": A.Morphological(scale=(7, 7), operation="dilation", p=0.5),
    # scale hyperparam
    "morpho_erosion": A.Morphological(scale=(7, 7), operation="erosion", p=0.5),
    "hflip": A.HorizontalFlip(p=0.5),
    "colorjitter": A.ColorJitter(
        brightness=(0.8, 1),
        contrast=(0.8, 1),
        saturation=(0.8, 1),
        hue=(-0.5, 0.5),
        always_apply=False,
        p=0.5,
    ),
    # blur_limit
    "gaussianblur": A.GaussianBlur(
        blur_limit=(3, 7), sigma_limit=0, always_apply=False, p=0.5
    ),
    # shift, scale, rotate
    "affine": A.ShiftScaleRotate(
        shift_limit=0.05,
        scale_limit=0.1,
        rotate_limit=30,
        border_mode=cv2.BORDER_CONSTANT,
        value=0,
        p=0.5,
    ),
    "invert": A.InvertImg(p=0.5),
    "gray": A.ToGray(p=0.5),
}
```

Composing data-augmentation given a string (Example: randomcrop224,colorjitter,invert,gray)

```
composed_transforms.append(transform_dict["resize256"])

# add all interim transforms
if transform_types != "":
    transform_types = transform_types.split(",")
    # Split the string into a list of transform types
    for transform_type in transform_types:
        if transform_type not in transform_dict:
            raise ValueError(f"Invalid transform_type: {transform_type}")
        composed_transforms.append(transform_dict[transform_type])
    else:
        assert False, "Not implemented!"

RESIZE_DIM = 96

# add all final preprocessing transforms
composed_transforms.extend(
    [
        # we resize all images to RESIZE_DIM
        A.Resize(height=RESIZE_DIM, width=RESIZE_DIM, p=1),
        A.Normalize(mean=(0.485, 0.456, 0.406), std=(0.229, 0.224, 0.225)),
        ToTensorV2(),
    ]
)

transform_type = A.Compose(composed_transforms)
test_or_val_transforms = A.Compose(
    [
        A.Resize(height=256, width=256, p=1),
        A.CenterCrop(height=224, width=224),

        A.Resize(height=RESIZE_DIM, width=RESIZE_DIM, p=1),
        A.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225]),
        ToTensorV2(),
    ]
)
```

Resizing all datapoints to 256

Performing data augmentations pipeline.(randomcrop224,colorjitter,invert,gray)

Standard resize to 96 to preserve memory

Primary 10 augmentations

Index	Augmentation	Type	Hyperparameters
1	invert	Pixel-Level	-
2	gray	Pixel-Level	-
3	gaussianblur	Pixel-Level	blur_limit = (3, 7), sigma_limit = 0
4	colorjitter	Pixel-Level	brightness = (0.8, 1), saturation = (0.8, 1), Contrast = (0.8, 1), hue = (-0.5, 0.5)
5	resize256	Spatial-Level	-
6	randomcrop256	Spatial-Level	-
7	hflip	Spatial-Level	-
8	morpho_erosion	Spatial-Level	kernel (w, h) = (7, 7)
9	morpho_dilation	Spatial-Level	kernel (w, h) = (7, 7)
10	affine	Spatial-Level	shift_limit = 0.05, scale_limit = 0.1, rotate_limit = 30

Ray Tune Hyperparameter Optimization Library

Ray tune hyper parameter optimization library

Provide list of hyper parameter

```
# ]
choices: [
    "resize256", "resize256,morpho_erosion", "resize256,morpho_dilation", "resize256,affine",
    "resize256,colorjitter", "resize256,hflip", "resize256,invert", "resize256,gaussianblur",
    "resize256,morpho_erosion,morpho_dilation", "resize256,morpho_erosion,affine",
    "resize256,morpho_erosion,colorjitter", "resize256,morpho_erosion,hflip", "resize256,morpho_erosion,invert",
    "resize256,morpho_erosion,gaussianblur", "resize256,morpho_dilation,affine", "resize256,morpho_dilation,colorjitter",
    "resize256,morpho_dilation,hflip", "resize256,morpho_dilation,invert", "resize256,morpho_dilation,gaussianblur",
    "resize256,affine,colorjitter", "resize256,affine,hflip", "resize256,affine,invert", "resize256,affine,gaussianblur",
    "resize256,colorjitter,hflip", "resize256,colorjitter,invert", "resize256,colorjitter,gaussianblur",
    "resize256,hflip,invert", "resize256,hflip,gaussianblur", "resize256,invert,gaussianblur",
    "resize256,morpho_erosion,morpho_dilation,affine", "resize256,morpho_erosion,morpho_dilation,colorjitter",
    "resize256,morpho_erosion,morpho_dilation,hflip", "resize256,morpho_erosion,morpho_dilation,invert",
    "resize256,morpho_erosion,morpho_dilation,gaussianblur", "resize256,morpho_erosion,affine,colorjitter",
    "resize256,morpho_erosion,affine,hflip", "resize256,morpho_erosion,affine,invert",
    "resize256,morpho_erosion,affine,gaussianblur", "resize256,morpho_erosion,colorjitter,hflip",
    "resize256,morpho_erosion,colorjitter,invert", "resize256,morpho_erosion,colorjitter,gaussianblur",
    "resize256,morpho_erosion,hflip,invert", "resize256,morpho_erosion,hflip,gaussianblur",
    "resize256,morpho_erosion,invert,gaussianblur", "resize256,morpho_dilation,affine,colorjitter",
    "resize256,morpho_dilation,affine,hflip", "resize256,morpho_dilation,affine,invert",
    "resize256,morpho_dilation,affine,gaussianblur", "resize256,morpho_dilation,colorjitter,hflip",
    "resize256,morpho_dilation,colorjitter,invert", "resize256,morpho_dilation,colorjitter,gaussianblur",
    "resize256,morpho_dilation,hflip,invert", "resize256,morpho_dilation,hflip,gaussianblur",
    "resize256,morpho_dilation,invert,gaussianblur", "resize256,affine,colorjitter,hflip",
    "resize256,affine,colorjitter,invert", "resize256,affine,colorjitter,gaussianblur",
    "resize256,affine,hflip,invert", "resize256,affine,hflip,gaussianblur",
    "resize256,affine,invert,gaussianblur", "resize256,colorjitter,hflip,invert",
    "resize256,colorjitter,hflip,gaussianblur", "resize256,colorjitter,invert,gaussianblur",
    "resize256,hflip,invert,gaussianblur"
]
```

We can use hyper band scheduler to efficiently search for best hyper parameter configurations. Eg: we can jointly search for different augmentations, different loss performance, different learning rates

93 combinations of augmentations

```
transform_type:
  choices: [
    "randomcrop224",

    "randomcrop224,morpho_erosion", "randomcrop224,morpho_dilation", "randomcrop224,affine", "randomcrop224,colorjitter",
    "randomcrop224,hflip", "randomcrop224,invert", "randomcrop224,gaussianblur", "randomcrop224,gray",

    "randomcrop224,morpho_erosion,morpho_dilation", "randomcrop224,morpho_erosion,affine", "randomcrop224,morpho_erosion,colorjitter",
    "randomcrop224,morpho_erosion,hflip", "randomcrop224,morpho_erosion,invert", "randomcrop224,morpho_erosion,gaussianblur",
    "randomcrop224,morpho_erosion,gray", "randomcrop224,morpho_dilation,affine", "randomcrop224,morpho_dilation,colorjitter",
    "randomcrop224,morpho_dilation,hflip", "randomcrop224,morpho_dilation,invert", "randomcrop224,morpho_dilation,gaussianblur",
    "randomcrop224,morpho_dilation,gray", "randomcrop224,affine,colorjitter", "randomcrop224,affine,hflip", "randomcrop224,affine,invert",
    "randomcrop224,affine,gaussianblur", "randomcrop224,affine,gray", "randomcrop224,colorjitter,hflip", "randomcrop224,colorjitter,invert",
    "randomcrop224,colorjitter,gaussianblur", "randomcrop224,colorjitter,gray", "randomcrop224,hflip,invert", "randomcrop224,hflip,gaussianblur",
    "randomcrop224,hflip,gray", "randomcrop224,invert,gaussianblur", "randomcrop224,invert,gray", "randomcrop224,gaussianblur,gray",

    "randomcrop224,morpho_erosion,morpho_dilation,affine", "randomcrop224,morpho_erosion,morpho_dilation,colorjitter",
    "randomcrop224,morpho_erosion,morpho_dilation,hflip", "randomcrop224,morpho_erosion,morpho_dilation,invert",
    "randomcrop224,morpho_erosion,morpho_dilation,gaussianblur", "randomcrop224,morpho_erosion,morpho_dilation,gray",
    "randomcrop224,morpho_erosion,affine,colorjitter", "randomcrop224,morpho_erosion,affine,hflip", "randomcrop224,morpho_erosion,affine,invert",
    "randomcrop224,morpho_erosion,affine,gaussianblur", "randomcrop224,morpho_erosion,affine,gray", "randomcrop224,morpho_erosion,colorjitter,hflip",
    "randomcrop224,morpho_erosion,colorjitter,invert", "randomcrop224,morpho_erosion,colorjitter,gaussianblur",
    "randomcrop224,morpho_erosion,colorjitter,gray", "randomcrop224,morpho_erosion,hflip,invert", "randomcrop224,morpho_erosion,hflip,gaussianblur",
    "randomcrop224,morpho_erosion,hflip,gray", "randomcrop224,morpho_erosion,invert,gaussianblur", "randomcrop224,morpho_erosion,invert,gray",
    "randomcrop224,morpho_erosion,gaussianblur,gray", "randomcrop224,morpho_dilation,affine,colorjitter", "randomcrop224,morpho_dilation,affine,hflip",
    "randomcrop224,morpho_dilation,affine,invert", "randomcrop224,morpho_dilation,affine,gaussianblur", "randomcrop224,morpho_dilation,affine,gray",
    "randomcrop224,morpho_dilation,colorjitter,hflip", "randomcrop224,morpho_dilation,colorjitter,invert",
    "randomcrop224,morpho_dilation,colorjitter,gaussianblur", "randomcrop224,morpho_dilation,colorjitter,gray", "randomcrop224,morpho_dilation,hflip,invert",
    "randomcrop224,morpho_dilation,hflip,gaussianblur", "randomcrop224,morpho_dilation,hflip,gray", "randomcrop224,morpho_dilation,invert,gaussianblur",
    "randomcrop224,morpho_dilation,invert,gray", "randomcrop224,morpho_dilation,gaussianblur,gray", "randomcrop224,affine,colorjitter,hflip",
    "randomcrop224,affine,colorjitter,invert", "randomcrop224,affine,colorjitter,gaussianblur", "randomcrop224,affine,colorjitter,gray",
    "randomcrop224,affine,hflip,invert", "randomcrop224,affine,hflip,gaussianblur", "randomcrop224,affine,hflip,gray", "randomcrop224,affine,invert,gaussianblur",
    "randomcrop224,affine,invert,gray", "randomcrop224,affine,gaussianblur,gray", "randomcrop224,colorjitter,hflip,invert", "randomcrop224,colorjitter,hflip,gaussianblur",
    "randomcrop224,colorjitter,hflip,gray", "randomcrop224,colorjitter,invert,gaussianblur", "randomcrop224,colorjitter,invert,gray", "randomcrop224,colorjitter,gaussianblur,gray",
    "randomcrop224,hflip,invert,gaussianblur", "randomcrop224,hflip,invert,gray", "randomcrop224,hflip,gaussianblur,gray", "randomcrop224,invert,gaussianblur,gray"
```

Results: without pretraining results on Alpub [ResNet18]

Experiment Settings

- Total cropped images: 34K
- Train-Validation-Test: 70%-15%-15%
- Models: **ResNet18**
- Optimizer: Adam with learning rate 0.001, 0.0003(SimCLR)
- Dataset: [ICDAR dataset](#)

Experiment	Dataset	Best Augmentation	Validation Acc.	Test Acc.
Baseline model	ICDAR	randomcrop224,morpho_erosion,morpho_dilation,gaussianblur	81.19%	80.67%
Triplet model	ICDAR	randomcrop224,morpho_dilation,affine, colorjitter	80.11%	79.16%
SimCLR model	ICDAR	randomcrop224,affine, colorjitter,gray	80.33%	80.00%

Results on ResNet-18 without pretraining on the Alpub dataset. We are presenting the best augmentations and corresponding validation and test accuracies, fine-tuned on ICDAR. The baseline model outperforms the other two methods.

Results: without pretraining results on Alpub [ResNet50]

Experiment Settings

- Total cropped images: 34K
- Train-Validation-Test: 70%-15%-15%
- Models: **ResNet50**
- Optimizer: Adam with learning rate 0.001, 0.0003(SimCLR)
- Dataset: [ICDAR dataset](#)

Experiment	Dataset	Best Augmentation	Validation Acc.	Test Ac.
Baseline model	ICDAR	randomcrop224,morpho_erosion, gaussianblur	80.70%	80.47%
Triplet model	ICDAR	randomcrop224,morpho_erosion, morpho_dilation,gaussianblur	79.29%	78.22%
SimCLR model	ICDAR	randomcrop224,colorjitter,gaussi anblur	80.05%	79.24%

Results on ResNet-50 without pretraining on the Alpub dataset

Selecting top-4 combinations

Strategy 1: T-test based selection

Augmentations	P-value
randomcrop224,morpho_dilation,hflip	0.008096
randomcrop224,hflip,gray	0.008721
randomcrop224,invert,gaussianblur,gray	0.008746
randomcrop224,colorjitter,hflip,invert	0.010912

Ran SimCLR with **3 runs** without pretraining using all 93 augmentations. Conducted a paired t-test comparing the base augmentation, RandomCrop224, against the other 92 augmentations. Below are the top-4 augmentations with their p-values. We compute the experiment on Alpub dataset using these 4 augmentations

Strategy 2: Best average validation accuracy

Augmentations	Average Test Acc
randomcrop224,morpho_erosion,morpho_dilation,affine	79.9413
randomcrop224,morpho_dilation,affine,colorjitter	79.8826
randomcrop224,morpho_erosion,affine,colorjitter	79.8434
randomcrop224,affine,colorjitter,gaussianblur	79.7521

Averaged the values of 3 runs of SimCLR without pretraining using all 93 augmentations. Then sorted the values. These are top-4 augmentations of averaged Test accuracy results.

Results: with pretraining results on Alpub [ResNet18] using Strategy 1

Experiment Settings

- Total cropped images: 205k
- Train-Validation-Test: 70%-15%-15%
- Models: **ResNet18**
- Optimizer: Adam with learning rate 0.001, 0.0003(SimCLR)
- Dataset: [Alpub dataset](#)

Experiment	Dataset	Best Augmentation	Validation Acc.	Test Acc.
Baseline model	Alpub + ICDAR	randomcrop224,hflip,gray	80.49%	79.94%
Triplet model	Alpub + ICDAR	randomcrop224,morpho_dilation,hflip	78.19%	77.51%
SimCLR model	Alpub + ICDAR	randomcrop224,colorjitter, hflip,invert	77.55%	76.14%

Results on ResNet-18 with pretraining on the Alpub dataset

Results: with pretraining results on Alpub [ResNet50] using Strategy 1

Experiment Settings

- Total cropped images: 205k
- Train-Validation-Test: 70%-15%-15%
- Models: **ResNet50**
- Optimizer: Adam with learning rate 0.001, 0.0003(SimCLR)
- Dataset: [Alpub dataset](#)

Experiment	Dataset	Best Augmentation	Validation Acc.	Test Acc.
Baseline model	Alpub + ICDAR	randomcrop224,morpho_dilation,hflip	80.21%	79.75%
Triplet model	Alpub + ICDAR	randomcrop224,invert,gaussianblur,gray	77.90%	77.03%
SimCLR model	Alpub + ICDAR	randomcrop224,invert,gaussianblur,gray	76.90%	76.59%

Results on ResNet-50 with pretraining on the Alpub dataset

Results: with pretraining results on Alpub [ResNet18] using Strategy 2

Experiment Settings

- Total cropped images: 205k
- Train-Validation-Test: 70%-15%-15%
- Models: **ResNet18**
- Optimizer: Adam with learning rate 0.001, 0.0003(SimCLR)
- Dataset: [Alpub dataset](#)

Experiment	Dataset	Best Augmentation	Validation Acc.	Test Ac.
Baseline model	Alpub + ICDAR	randomcrop224,morpho_erosion,affine,colorjitter	80.68%	81.14%
Triplet model	Alpub + ICDAR	randomcrop224,morpho_dilation,affine,colorjitter	79.57%	78.88%
SimCLR model	Alpub + ICDAR	randomcrop224,morpho_erosion, affine,colorjitter	79.94%	79.18%

Results on ResNet-18 with pretraining on the Alpub dataset

Results: with pretraining results on Alpub [ResNet50] using Strategy 2

Experiment Settings

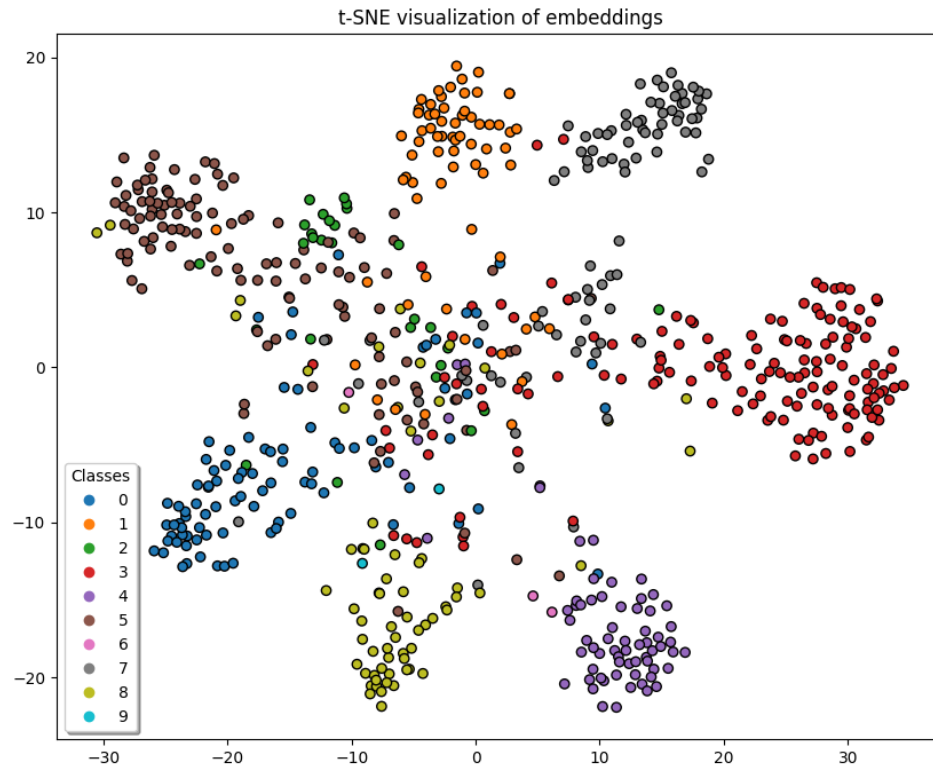
- Total cropped images: 205k
- Train-Validation-Test: 70%-15%-15%
- Models: **ResNet50**
- Optimizer: Adam with learning rate 0.001, 0.0003(SimCLR)
- Dataset: [Alpub dataset](#)

Experiment	Dataset	Best Augmentation	Validation Acc.	Test Ac.
Baseline model	Alpub + ICDAR	randomcrop224, affine, colorjitter, gaussianblur	81.35%	81.17%
Triplet model	Alpub + ICDAR	randomcrop224,morpho_dilation,affine, colorjitter	79.17%	78.24%
SimCLR model	Alpub + ICDAR	randomcrop224, affine, colorjitter, gaussianblur	78.68%	78.85%

Results on ResNet-50 with pretraining on the Alpub dataset

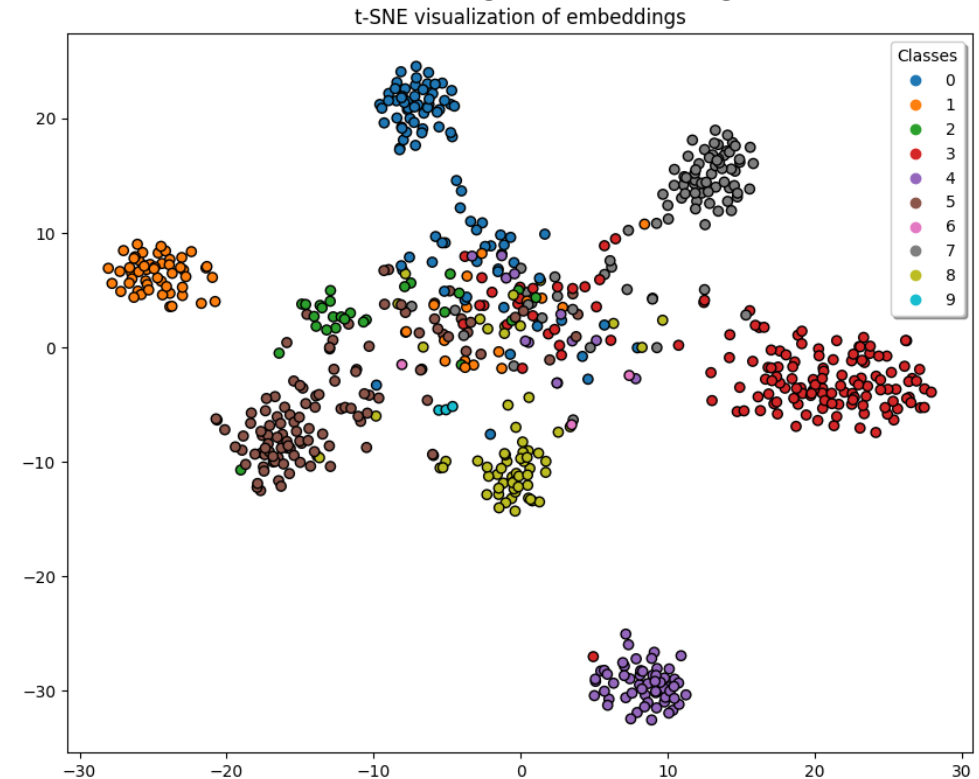
tSNE embeddings visualization of Baseline model

After pretraining embeddings



Embeddings of ICDAR dataset images after pretraining the model on Alpub dataset for 20 epochs with CE loss on ResNet-18

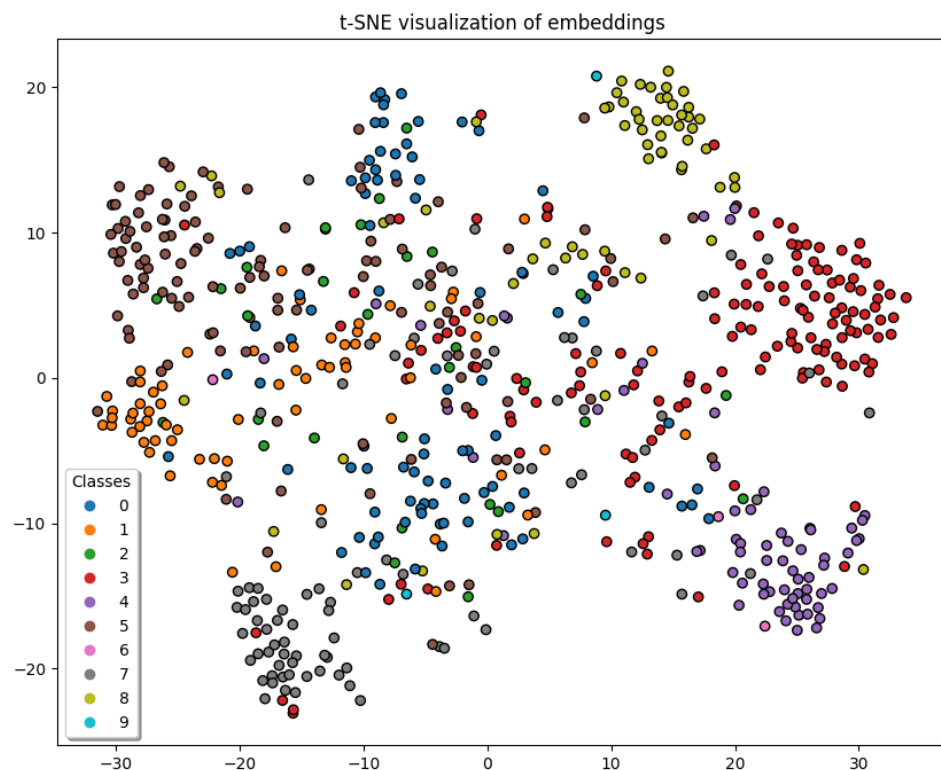
After finetuning embeddings



Embeddings of ICDAR dataset images after the model further finetuned on all layers for 20 epochs with CE loss on ResNet-18

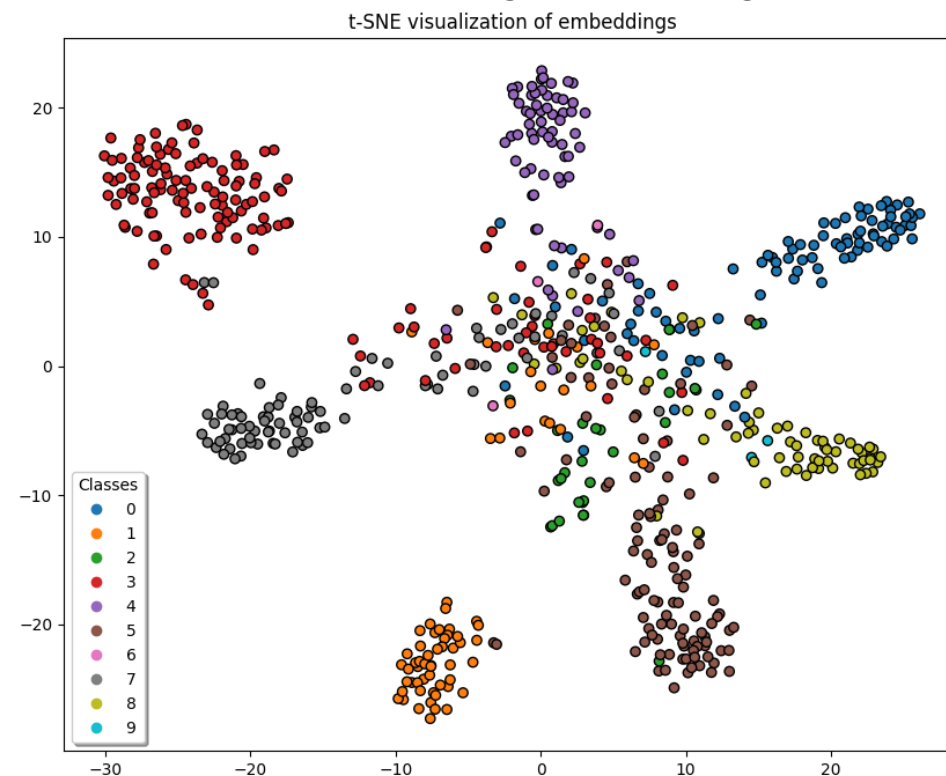
tSNE embeddings visualization of Triplet Model

After pretraining embeddings



Embeddings of ICDAR dataset images after pretraining the model on Alpub for **20 epochs** with Triplet loss on ResNet-18


After finetuning embeddings



Embeddings of ICDAR dataset images after the model further **finetuned** on all layers for **20 epochs** with CE loss on ResNet-18

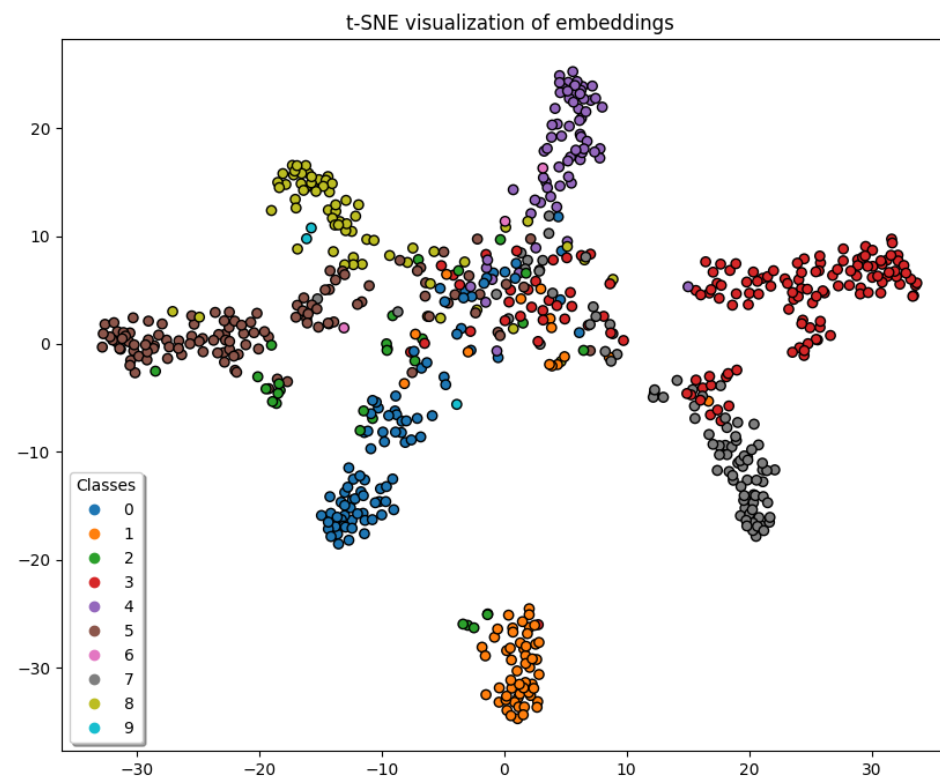
tSNE embeddings visualization of SimCLR model

After pretraining embeddings



Embeddings of ICDAR dataset images after pretraining the model on Alpub dataset for 20 epochs with InfoNCE loss on ResNet-18

After finetuning embeddings



Embeddings of ICDAR dataset images after the model further finetuned on all layers for 20 epochs with CE loss on ResNet-18

Discussion and Limitations



- ☐ Hyper Parameter Tuning
- ☐ Data augmentation
- ☐ Cropping size of SimCLR
- ☐ Batch size in SimCLR
- ☐ Dataset construction

SimCLR Cropping Leads to Semantic Shift. We observe the two views of the image cropped from the original image with 60% area. It can be seen that this cropping scheme leads to a change in the labels.

Conclusion and Future Work

- SimCLR not performing well than the traditional methods, cross-entropy and triplet loss, for letter recognition.
- Using different data augmentations, especially cropping, caused SimCLR to struggle due to changes in the meaning of the images.
- The traditional cross-entropy model consistently performed better than SimCLR and triplet loss models.
- t-SNE plots show that SimCLR failed to form well-defined clusters during pretraining.

Future Work:

- Explore alternative contrastive learning strategies.
- Tune hyperparameters and refine data augmentation techniques for improved performance and other limitations

THANK YOU